

5 Flower Types Classification Dataset

The **5 Flower Types Classification Dataset** is a collection of images belonging to five different flower classes: Lilly, Lotus, Sunflower, Orchid, and Tulip. Each flower class contains 1000 images, resulting in a total of 5000 images in the dataset.

This dataset is suitable for training and evaluating a multi-class Convolutional Neural Network (CNN) model to classify flower images into one of the five mentioned classes. The goal of the classification task is to accurately identify the type of flower from an input image.

The dataset can be used to explore various deep learning techniques for image classification, such as data augmentation, transfer learning, and model fine-tuning. It provides a challenging task due to the visual similarity and subtle differences among different flower types.

Dataset Details:

- Number of classes: 5
- Total images: 5000 (1000 images per class)
- Image format: JPG or PNG
- Image resolution: Varies (please preprocess the images to a consistent size if required)

The 5 Flower Types Classification Dataset is a valuable resource for researchers, students, and practitioners interested in the field of computer vision, specifically in image classification tasks. It can be used for educational purposes, benchmarking different models, and advancing the state-of-the-art in flower classification.

Feel free to download the dataset and start exploring the fascinating world of flower image classification!:

<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>

<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>



The Essence of Transfer Learning

- Transfer learning, a pivotal concept in AI, empowers models to leverage knowledge from one task and apply it to another. By drawing inspiration from human learning, this approach has redefined machine learning. Instead of starting from scratch for each task, models are pre-trained on vast datasets and

then fine-tuned for specific applications, saving time and resources.

- In the realm of computer vision, advanced models such as ResNet exhibit exceptional proficiency in image comprehension, enabling them to excel in tasks like medical image analysis and autonomous vehicle navigation. However, the field faces its own set of challenges, including issues like negative transfer and bias amplification, underscoring the importance of careful and considerate implementation strategies.
- The future of transfer learning is bright, with ongoing research focusing on enhancing adaptability and reducing fine-tuning requirements. As AI continues to evolve, transfer learning remains a cornerstone, propelling us toward a future where machines seamlessly learn, adapt, and revolutionize industries.

```
In [1]: 1 # import libraries
2 import os
3 import shutil
4 import random
5 import numpy as np
6 import pandas as pd
7 import datetime
8 from matplotlib import pyplot as plt
9 from PIL import Image
10 from sklearn.metrics import (accuracy_score,
11                             classification_report,
12                             confusion_matrix)
13 import tensorflow as tf
14 from tensorflow.keras.preprocessing.image import ImageDataGenerator
15 from tensorflow.keras.optimizers import Adam
16 from tensorflow.keras.callbacks import EarlyStopping
17 from tensorflow.keras.layers import (Conv2D, MaxPool2D,
18                                     Flatten, Dense, Dropout)
19 from tensorflow.keras import Sequential
20 import tensorflow_hub as hub
21 os.environ['KAGGLE_CONFIG_DIR'] = '/content'
22 BATCH_SIZE = 32
23 LEARNING_RATE = 0.0001
24 TARGET_SIZE = (256, 256)
```

Downloading and preparing data for model

```
In [2]: 1 !kaggle datasets download -d kausthubkannan/5-flower-types-classification-dataset
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /content/kaggle.json'
Downloading 5-flower-types-classification-dataset.zip to /content
 97% 235M/242M [00:02<00:00, 71.9MB/s]
100% 242M/242M [00:03<00:00, 82.2MB/s]
```

```
In [3]: 1 !unzip \*.zip && rm *.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: flower_images/Lilly/00048a5c76.jpg
inflating: flower_images/Lilly/001ff6644e.jpg
inflating: flower_images/Lilly/001ff6656j.jpg
inflating: flower_images/Lilly/00973ad1b1.jpg
inflating: flower_images/Lilly/00a7d512d6.jpg
inflating: flower_images/Lilly/00f36a3c40.jpg
inflating: flower_images/Lilly/013628cccc.jpg
inflating: flower_images/Lilly/01998d6fb5.jpg
inflating: flower_images/Lilly/01a0ec319c.jpg
inflating: flower_images/Lilly/01b4bb0289.jpg
inflating: flower_images/Lilly/025ef3ea44.jpg
inflating: flower_images/Lilly/02a7a2df46.jpg
inflating: flower_images/Lilly/02be2ca388.jpg
inflating: flower_images/Lilly/035cce082f.jpg
inflating: flower_images/Lilly/039eba79d4.jpg
inflating: flower_images/Lilly/04067b91d6.jpg
inflating: flower_images/Lilly/04acfd5449.jpg
inflating: flower_images/Lilly/05777790e2.jpg
```

```
In [4]: 1 !pip install split-folders
```

Collecting split-folders

Downloading split_folders-0.5.1-py3-none-any.whl (8.4 kB)

Installing collected packages: split-folders

Successfully installed split-folders-0.5.1

```
In [5]: 1 import splitfolders
```

```
In [6]: 1 src_dir = '/content/flower_images'
2 dst_dir = '/content/Data'
```

Data preprocessing (image augmentation)

```
In [7]: 1 splitfolders.ratio(input=src_dir, output=dst_dir, ratio=(0.6, 0.2, 0.2))
```

Copying files: 5000 files [00:01, 4643.77 files/s]

```
In [8]: 1 for dpath, dname, filename in os.walk("Data"):
2         print(f"looking at {dpath} found {dname} folders files available {len(filename)}")
```

looking at Data found ['train', 'val', 'test'] folders files available 0
looking at Data/train found ['Sunflower', 'Orchid', 'Tulip', 'Lilly', 'Lotus'] folders
files available 0
looking at Data/train/Sunflower found [] folders files available 600
looking at Data/train/Orchid found [] folders files available 600
looking at Data/train/Tulip found [] folders files available 600
looking at Data/train/Lilly found [] folders files available 600
looking at Data/train/Lotus found [] folders files available 600
looking at Data/val found ['Sunflower', 'Orchid', 'Tulip', 'Lilly', 'Lotus'] folders fi
les available 0
looking at Data/val/Sunflower found [] folders files available 200
looking at Data/val/Orchid found [] folders files available 200
looking at Data/val/Tulip found [] folders files available 200
looking at Data/val/Lilly found [] folders files available 200
looking at Data/val/Lotus found [] folders files available 200
looking at Data/test found ['Sunflower', 'Orchid', 'Tulip', 'Lilly', 'Lotus'] folders f
iles available 0
looking at Data/test/Sunflower found [] folders files available 200
looking at Data/test/Orchid found [] folders files available 200
looking at Data/test/Tulip found [] folders files available 200
looking at Data/test/Lilly found [] folders files available 200
looking at Data/test/Lotus found [] folders files available 200







```
In [9]: 1 print("Training Images.")
2 train_datagen = ImageDataGenerator(rescale=1/255., rotation_range=0.2,
3                                     # brightness_range=(0.2, 0.5),
4                                     zoom_range=0.2, shear_range=0.2,
5                                     horizontal_flip=True)
6 train_dataset = train_datagen.flow_from_directory('/content/Data/train',
7                                                    target_size=TARGET_SIZE,
8                                                    batch_size=BATCH_SIZE,
9                                                    shuffle=True)
10 print("Validation Images")
11 val_datagen = ImageDataGenerator(rescale=1/255.)
12 val_dataset = val_datagen.flow_from_directory('/content/Data/val',
13                                                target_size=TARGET_SIZE,
14                                                batch_size=BATCH_SIZE,
15                                                shuffle=False)
16
17 print("Testing Images.")
18 test_datagen = ImageDataGenerator(rescale=1/255.)
19 test_dataset = test_datagen.flow_from_directory('/content/Data/test',
20                                                  target_size=TARGET_SIZE,
21                                                  batch_size=BATCH_SIZE,
22                                                  shuffle=False)
```

Training Images.
Found 3000 images belonging to 5 classes.
Validation Images
Found 1000 images belonging to 5 classes.
Testing Images.
Found 1000 images belonging to 5 classes.

```

In [10]: 1 images, labels = next(train_dataset)
          2 labels = np.argmax(labels, axis=1)
          3 class_names = list(train_dataset.class_indices.keys())
          4 def plot_random_images(images, labels, class_names):
          5     plt.figure(figsize=(12, 6))
          6
          7     for i in range(6):
          8         ax = plt.subplot(2, 3, i+1)
          9         rand_index = random.choice(range(len(images)))
         10         plt.imshow(images[rand_index])
         11         plt.title(class_names[labels[rand_index]], color='green', fontsize=12)
         12
         13     plt.tight_layout()
         14     plt.show()
         15
         16 plot_random_images(images, labels, class_names)

```



Baseline Model

```
In [11]: 1 baseline_model = Sequential([
2
3         Conv2D(filters=16, kernel_size=(3,3), strides=1,
4             activation='relu', input_shape=(256, 256, 3)),
5         MaxPool2D(pool_size=(2,2), strides=2, padding='valid'),
6
7         Conv2D(filters=32, kernel_size=(3,3), strides=2,
8             activation='relu'),
9         MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
10
11        Conv2D(filters=64, kernel_size=(3,3), strides=2,
12            activation='relu'),
13        MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
14
15        Flatten(),
16        Dense(256, activation='relu'),
17        Dense(128, activation='relu'),
18        Dense(64, activation='relu'),
19        Dense(5, activation='softmax')
20
21    ])
```

```
In [12]: 1 early_stop = EarlyStopping(monitor="val_loss", patience=5)
```

```
In [13]: 1 baseline_model.compile(optimizer=Adam(learning_rate=LEARNING_RATE),
2         loss=tf.keras.losses.CategoricalCrossentropy(),
3         metrics=['accuracy'])
```

```
In [14]: 1 baseline_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 16)	448
max_pooling2d (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_1 (Conv2D)	(None, 63, 63, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_2 (Conv2D)	(None, 31, 31, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 64)	0
flatten (Flatten)	(None, 61504)	0
dense (Dense)	(None, 256)	15745280
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 5)	325
=====		
Total params: 15,810,341		
Trainable params: 15,810,341		
Non-trainable params: 0		

[illegible]

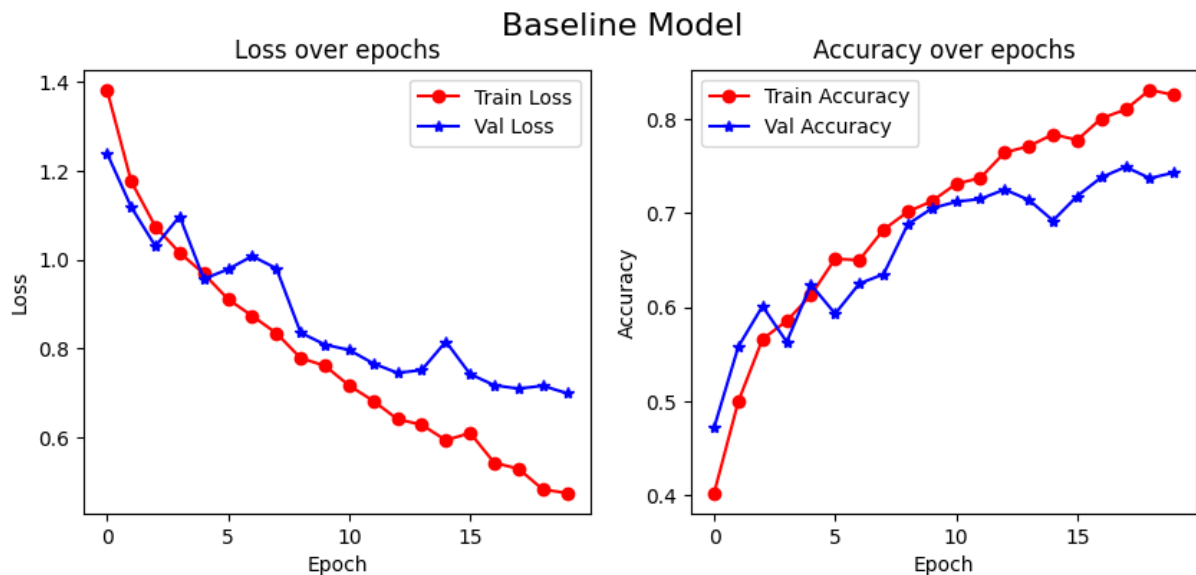
Epoch 1/20
94/94 [=====] - 67s 563ms/step - loss: 1.3815 - accuracy: 0.40
23 - val_loss: 1.2391 - val_accuracy: 0.4720
Epoch 2/20
94/94 [=====] - 54s 572ms/step - loss: 1.1757 - accuracy: 0.49
97 - val_loss: 1.1182 - val_accuracy: 0.5580
Epoch 3/20
94/94 [=====] - 52s 557ms/step - loss: 1.0739 - accuracy: 0.56
60 - val_loss: 1.0302 - val_accuracy: 0.6010
Epoch 4/20
94/94 [=====] - 51s 544ms/step - loss: 1.0156 - accuracy: 0.58
53 - val_loss: 1.0972 - val_accuracy: 0.5630
Epoch 5/20
94/94 [=====] - 51s 544ms/step - loss: 0.9681 - accuracy: 0.61
27 - val_loss: 0.9564 - val_accuracy: 0.6240
Epoch 6/20
94/94 [=====] - 51s 546ms/step - loss: 0.9110 - accuracy: 0.65
13 - val_loss: 0.9783 - val_accuracy: 0.5930
Epoch 7/20
94/94 [=====] - 52s 555ms/step - loss: 0.8733 - accuracy: 0.65
00 - val_loss: 1.0077 - val_accuracy: 0.6250
Epoch 8/20
94/94 [=====] - 53s 564ms/step - loss: 0.8346 - accuracy: 0.68
23 - val_loss: 0.9801 - val_accuracy: 0.6350
Epoch 9/20
94/94 [=====] - 51s 541ms/step - loss: 0.7787 - accuracy: 0.70
13 - val_loss: 0.8357 - val_accuracy: 0.6880
Epoch 10/20
94/94 [=====] - 51s 542ms/step - loss: 0.7615 - accuracy: 0.71
27 - val_loss: 0.8090 - val_accuracy: 0.7050
Epoch 11/20
94/94 [=====] - 52s 548ms/step - loss: 0.7171 - accuracy: 0.73
10 - val_loss: 0.7969 - val_accuracy: 0.7120
Epoch 12/20
94/94 [=====] - 51s 547ms/step - loss: 0.6826 - accuracy: 0.73
73 - val_loss: 0.7668 - val_accuracy: 0.7150
Epoch 13/20
94/94 [=====] - 52s 553ms/step - loss: 0.6414 - accuracy: 0.76
43 - val_loss: 0.7453 - val_accuracy: 0.7250
Epoch 14/20
94/94 [=====] - 52s 559ms/step - loss: 0.6289 - accuracy: 0.77
10 - val_loss: 0.7518 - val_accuracy: 0.7140
Epoch 15/20
94/94 [=====] - 52s 551ms/step - loss: 0.5940 - accuracy: 0.78
37 - val_loss: 0.8158 - val_accuracy: 0.6920
Epoch 16/20
94/94 [=====] - 54s 572ms/step - loss: 0.6105 - accuracy: 0.77
77 - val_loss: 0.7427 - val_accuracy: 0.7180
Epoch 17/20
94/94 [=====] - 51s 540ms/step - loss: 0.5433 - accuracy: 0.80
07 - val_loss: 0.7176 - val_accuracy: 0.7380
Epoch 18/20
94/94 [=====] - 52s 554ms/step - loss: 0.5301 - accuracy: 0.81
00 - val_loss: 0.7099 - val_accuracy: 0.7490
Epoch 19/20
94/94 [=====] - 50s 537ms/step - loss: 0.4837 - accuracy: 0.83
10 - val_loss: 0.7165 - val_accuracy: 0.7370
Epoch 20/20
94/94 [=====] - 51s 538ms/step - loss: 0.4754 - accuracy: 0.82
53 - val_loss: 0.6996 - val_accuracy: 0.7430

```

In [16]: 1 def plot_predictions(loss_df, title):
2         """
3         Plots the training and validation loss, as well as training and validation
4         accuracy, over epochs for a given dataset.
5
6         Parameters:
7         loss_df (pandas.DataFrame): A DataFrame containing loss and accuracy values
8                                     for both training and validation sets over
9                                     epochs. Columns should include 'loss', 'val_loss'
10                                    'accuracy', and 'val_accuracy'.
11         title (str): Title for the entire plot.
12
13         Returns:
14         None (Displays the plot with loss and accuracy curves)
15         """
16         fig, ax = plt.subplots(1, 2, figsize=(10, 4))
17
18         ax[0].plot(loss_df['loss'], color='red', marker='o',
19                   label='Train Loss')
20         ax[0].plot(loss_df['val_loss'], color='blue', marker='*',
21                   label='Val Loss')
22
23         ax[0].set_title('Loss over epochs')
24         ax[0].set_xlabel('Epoch')
25         ax[0].set_ylabel('Loss')
26         ax[0].legend()
27
28         ax[1].plot(loss_df['accuracy'], color='red', marker='o',
29                   label='Train Accuracy')
30         ax[1].plot(loss_df['val_accuracy'], color='blue', marker='*',
31                   label='Val Accuracy')
32
33         ax[1].set_title('Accuracy over epochs')
34         ax[1].set_xlabel('Epoch')
35         ax[1].set_ylabel('Accuracy')
36         ax[1].legend()
37
38         fig.suptitle(title, fontsize=16)
39         plt.show()

```

```
In [17]: 1 loss_df = pd.DataFrame(baseline_model_history.history)
2         plot_predictions(loss_df, title="Baseline Model")
```



```
In [18]: 1 def create_tensorboard_callback(dir_name, experiment_name):
2         """
3         Creates a TensorBoard callback for tracking training progress and
4         visualizing metrics.
5
6         Parameters:
7         dir_name (str): The base directory where TensorBoard logs will be stored.
8         experiment_name (str): A unique name for the experiment, used to create a
9         subdirectory within dir_name.
10
11         Returns:
12         tf.keras.callbacks.TensorBoard: A TensorBoard callback instance configured
13         with the appropriate log directory.
14         """
15         log_dir = (dir_name + "/" + experiment_name + "/" +
16                   datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
17         tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
18         return tensorboard_callback
19
```

```
In [19]: 1 resnet_url = "https://tfhub.dev/google/imagenet/resnet_v2_50/feature_vector/5"
2         efficientnet_url = "https://tfhub.dev/tensorflow/efficientnet/b0/feature-vector/1"
```

```
In [20]: 1 def create_model(model_url, num_classes):
2         """
3         Creates a Keras model for transfer learning using a pre-trained feature
4         extraction layer from TensorFlow Hub.
5
6         Parameters:
7             model_url (str): The URL of the pre-trained model from TensorFlow Hub that
8             will be used as the feature extraction layer. This model should be
9             compatible with KerasLayer.
10            num_classes (int, optional): The number of classes for the final
11            classification layer. Default is 5.
12            TARGET_SIZE (tuple, optional): The target input image size expected by the
13            model's feature extraction layer. It's represented as a tuple of two
14            integers (height, width). Default is (256, 256).
15
16        Returns:
17            A Keras Sequential model with the following structure:
18                1. A KerasLayer with the pre-trained feature extraction model from
19                TensorFlow Hub.
20                2. A Dense layer with softmax activation for final classification.
21
22        """
23        feature_extraction_layer = hub.KerasLayer(model_url, trainable=False,
24                                                  name="feature_extraction_layer",
25                                                  input_shape=TARGET_SIZE + (3,))
26        model = Sequential([
27            feature_extraction_layer,
28            Dense(num_classes, activation="softmax", name='output_ayer')
29        ])
30        return model
```

```
In [21]: 1 resnet_model = create_model(resnet_url, num_classes=train_dataset.num_classes)
```

```
In [22]: 1 resnet_model.summary()
```

Model: "sequential_1"

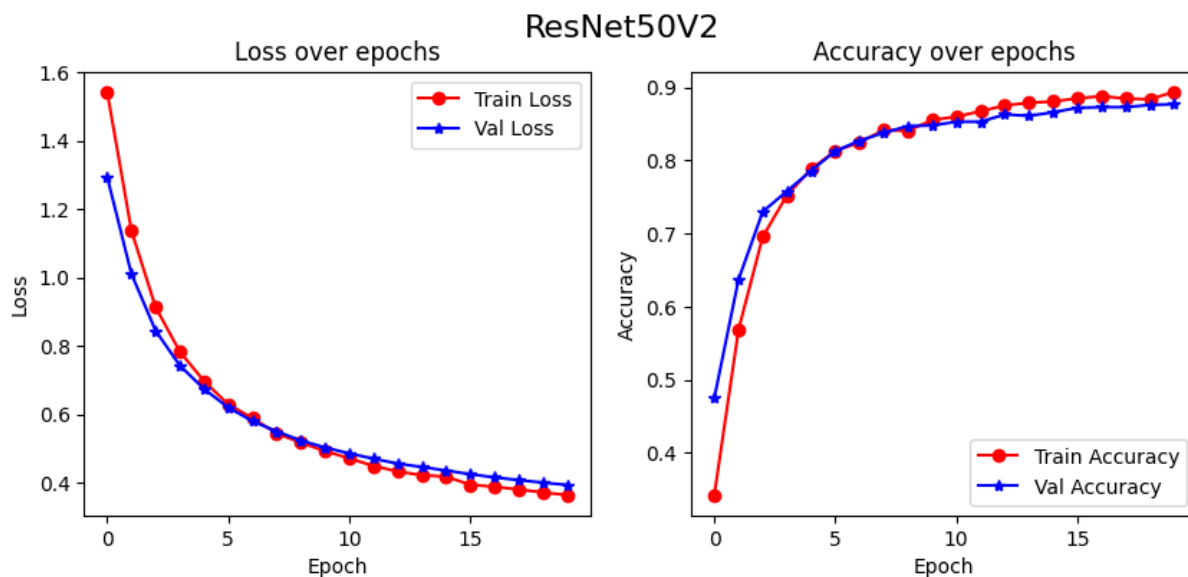
Layer (type)	Output Shape	Param #
=====		
feature_extraction_layer (KerasLayer)	(None, 2048)	23564800
output_ayer (Dense)	(None, 5)	10245
=====		
Total params: 23,575,045		
Trainable params: 10,245		
Non-trainable params: 23,564,800		
=====		

```
In [23]: 1 resnet_model.compile(loss=tf.keras.losses.CategoricalCrossentropy(),
2                             optimizer=tf.keras.optimizers.Adam(learning_rate=LEARNING_RATE)
3                             metrics=["accuracy"])
```

```
In [25]: 1 resnet_history = resnet_model.fit(train_dataset, epochs=20,  
2                                           callbacks = [create_tensorboard_callback(  
3                                           dir_name="tensorflow_hub",  
4                                           experiment_name="Resnet50V2"  
5                                           )],  
6                                           validation_data=val_dataset  
7                                           )
```

Epoch 1/20
94/94 [=====] - 64s 607ms/step - loss: 1.5424 - accuracy: 0.34
23 - val_loss: 1.2940 - val_accuracy: 0.4750
Epoch 2/20
94/94 [=====] - 56s 593ms/step - loss: 1.1369 - accuracy: 0.56
83 - val_loss: 1.0104 - val_accuracy: 0.6370
Epoch 3/20
94/94 [=====] - 55s 582ms/step - loss: 0.9127 - accuracy: 0.69
57 - val_loss: 0.8442 - val_accuracy: 0.7300
Epoch 4/20
94/94 [=====] - 54s 571ms/step - loss: 0.7827 - accuracy: 0.75
23 - val_loss: 0.7422 - val_accuracy: 0.7580
Epoch 5/20
94/94 [=====] - 55s 587ms/step - loss: 0.6969 - accuracy: 0.78
83 - val_loss: 0.6738 - val_accuracy: 0.7860
Epoch 6/20
94/94 [=====] - 55s 588ms/step - loss: 0.6287 - accuracy: 0.81
30 - val_loss: 0.6197 - val_accuracy: 0.8130
Epoch 7/20
94/94 [=====] - 54s 568ms/step - loss: 0.5887 - accuracy: 0.82
47 - val_loss: 0.5797 - val_accuracy: 0.8270
Epoch 8/20
94/94 [=====] - 55s 589ms/step - loss: 0.5437 - accuracy: 0.84
17 - val_loss: 0.5496 - val_accuracy: 0.8380
Epoch 9/20
94/94 [=====] - 54s 572ms/step - loss: 0.5168 - accuracy: 0.84
07 - val_loss: 0.5230 - val_accuracy: 0.8470
Epoch 10/20
94/94 [=====] - 54s 571ms/step - loss: 0.4920 - accuracy: 0.85
53 - val_loss: 0.5026 - val_accuracy: 0.8480
Epoch 11/20
94/94 [=====] - 57s 602ms/step - loss: 0.4710 - accuracy: 0.86
00 - val_loss: 0.4853 - val_accuracy: 0.8530
Epoch 12/20
94/94 [=====] - 53s 568ms/step - loss: 0.4486 - accuracy: 0.86
73 - val_loss: 0.4699 - val_accuracy: 0.8530
Epoch 13/20
94/94 [=====] - 53s 569ms/step - loss: 0.4326 - accuracy: 0.87
53 - val_loss: 0.4555 - val_accuracy: 0.8630
Epoch 14/20
94/94 [=====] - 55s 585ms/step - loss: 0.4215 - accuracy: 0.87
90 - val_loss: 0.4456 - val_accuracy: 0.8610
Epoch 15/20
94/94 [=====] - 54s 579ms/step - loss: 0.4176 - accuracy: 0.88
07 - val_loss: 0.4352 - val_accuracy: 0.8660
Epoch 16/20
94/94 [=====] - 54s 568ms/step - loss: 0.3937 - accuracy: 0.88
50 - val_loss: 0.4245 - val_accuracy: 0.8720
Epoch 17/20
94/94 [=====] - 54s 570ms/step - loss: 0.3879 - accuracy: 0.88
80 - val_loss: 0.4153 - val_accuracy: 0.8730
Epoch 18/20
94/94 [=====] - 54s 576ms/step - loss: 0.3797 - accuracy: 0.88
47 - val_loss: 0.4073 - val_accuracy: 0.8730
Epoch 19/20
94/94 [=====] - 53s 562ms/step - loss: 0.3715 - accuracy: 0.88
37 - val_loss: 0.3999 - val_accuracy: 0.8760
Epoch 20/20
94/94 [=====] - 53s 567ms/step - loss: 0.3635 - accuracy: 0.89
40 - val_loss: 0.3932 - val_accuracy: 0.8770

```
In [26]: 1 loss_df = pd.DataFrame(resnet_history.history)
2         plot_predictions(loss_df, title="ResNet50V2")
```



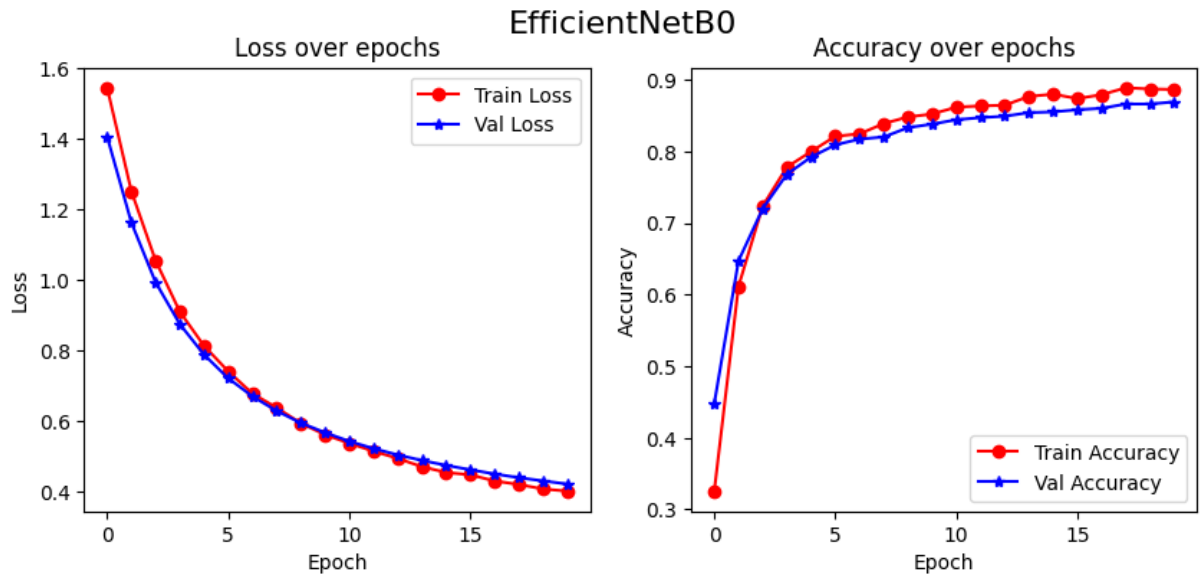
```
In [27]: 1 efficientnet_model = create_model(model_url=efficientnet_url,
2                                           num_classes=train_dataset.num_classes)
```

```
In [28]: 1 efficientnet_model.compile(loss=tf.keras.losses.CategoricalCrossentropy(),
2                                optimizer=tf.keras.optimizers.Adam(learning_rate=LEARNING_RATE)
3                                metrics=["accuracy"])
```

```
In [29]: 1 effecientnet_history = efficientnet_model.fit(train_dataset, epochs=20,  
2                                                    callbacks = [create_tensorboard_callback(  
3                                                         dir_name="tensorflow_hub",  
4                                                         experiment_name="Effecientnet"  
5                                                         )],  
6                                                    validation_data=val_dataset  
7                                                    )
```


Epoch 1/20
94/94 [=====] - 67s 588ms/step - loss: 1.5451 - accuracy: 0.32
57 - val_loss: 1.4065 - val_accuracy: 0.4480
Epoch 2/20
94/94 [=====] - 54s 574ms/step - loss: 1.2521 - accuracy: 0.61
13 - val_loss: 1.1632 - val_accuracy: 0.6470
Epoch 3/20
94/94 [=====] - 53s 562ms/step - loss: 1.0557 - accuracy: 0.72
37 - val_loss: 0.9943 - val_accuracy: 0.7200
Epoch 4/20
94/94 [=====] - 54s 579ms/step - loss: 0.9111 - accuracy: 0.77
80 - val_loss: 0.8757 - val_accuracy: 0.7680
Epoch 5/20
94/94 [=====] - 53s 561ms/step - loss: 0.8144 - accuracy: 0.79
97 - val_loss: 0.7893 - val_accuracy: 0.7920
Epoch 6/20
94/94 [=====] - 54s 572ms/step - loss: 0.7420 - accuracy: 0.82
10 - val_loss: 0.7225 - val_accuracy: 0.8090
Epoch 7/20
94/94 [=====] - 52s 551ms/step - loss: 0.6791 - accuracy: 0.82
43 - val_loss: 0.6712 - val_accuracy: 0.8170
Epoch 8/20
94/94 [=====] - 53s 567ms/step - loss: 0.6391 - accuracy: 0.83
87 - val_loss: 0.6299 - val_accuracy: 0.8200
Epoch 9/20
94/94 [=====] - 53s 564ms/step - loss: 0.5950 - accuracy: 0.84
80 - val_loss: 0.5958 - val_accuracy: 0.8330
Epoch 10/20
94/94 [=====] - 52s 556ms/step - loss: 0.5624 - accuracy: 0.85
20 - val_loss: 0.5686 - val_accuracy: 0.8380
Epoch 11/20
94/94 [=====] - 61s 651ms/step - loss: 0.5382 - accuracy: 0.86
17 - val_loss: 0.5439 - val_accuracy: 0.8440
Epoch 12/20
94/94 [=====] - 57s 605ms/step - loss: 0.5156 - accuracy: 0.86
33 - val_loss: 0.5234 - val_accuracy: 0.8470
Epoch 13/20
94/94 [=====] - 55s 575ms/step - loss: 0.4954 - accuracy: 0.86
43 - val_loss: 0.5055 - val_accuracy: 0.8490
Epoch 14/20
94/94 [=====] - 53s 561ms/step - loss: 0.4718 - accuracy: 0.87
67 - val_loss: 0.4896 - val_accuracy: 0.8540
Epoch 15/20
94/94 [=====] - 55s 583ms/step - loss: 0.4556 - accuracy: 0.87
97 - val_loss: 0.4759 - val_accuracy: 0.8550
Epoch 16/20
94/94 [=====] - 53s 567ms/step - loss: 0.4489 - accuracy: 0.87
37 - val_loss: 0.4632 - val_accuracy: 0.8580
Epoch 17/20
94/94 [=====] - 53s 567ms/step - loss: 0.4316 - accuracy: 0.87
90 - val_loss: 0.4515 - val_accuracy: 0.8600
Epoch 18/20
94/94 [=====] - 53s 564ms/step - loss: 0.4216 - accuracy: 0.88
87 - val_loss: 0.4409 - val_accuracy: 0.8660
Epoch 19/20
94/94 [=====] - 53s 562ms/step - loss: 0.4088 - accuracy: 0.88
70 - val_loss: 0.4317 - val_accuracy: 0.8660
Epoch 20/20
94/94 [=====] - 54s 573ms/step - loss: 0.4028 - accuracy: 0.88
63 - val_loss: 0.4229 - val_accuracy: 0.8690

```
In [30]: 1 loss_df = pd.DataFrame(efficientnet_history.history)
2         plot_predictions(loss_df, title="EfficientNetB0")
```



making predictions

```
In [41]: 1 baseline_prediction = np.argmax(baseline_model.predict(test_dataset),
2                                       axis=1)
3         resnet_prediction = np.argmax(resnet_model.predict(test_dataset),
4                                       axis=1)
5         efficientnet_prediction = np.argmax(efficientnet_model.predict(test_dataset),
6                                           axis=1)
7         y_true = test_dataset.labels
```

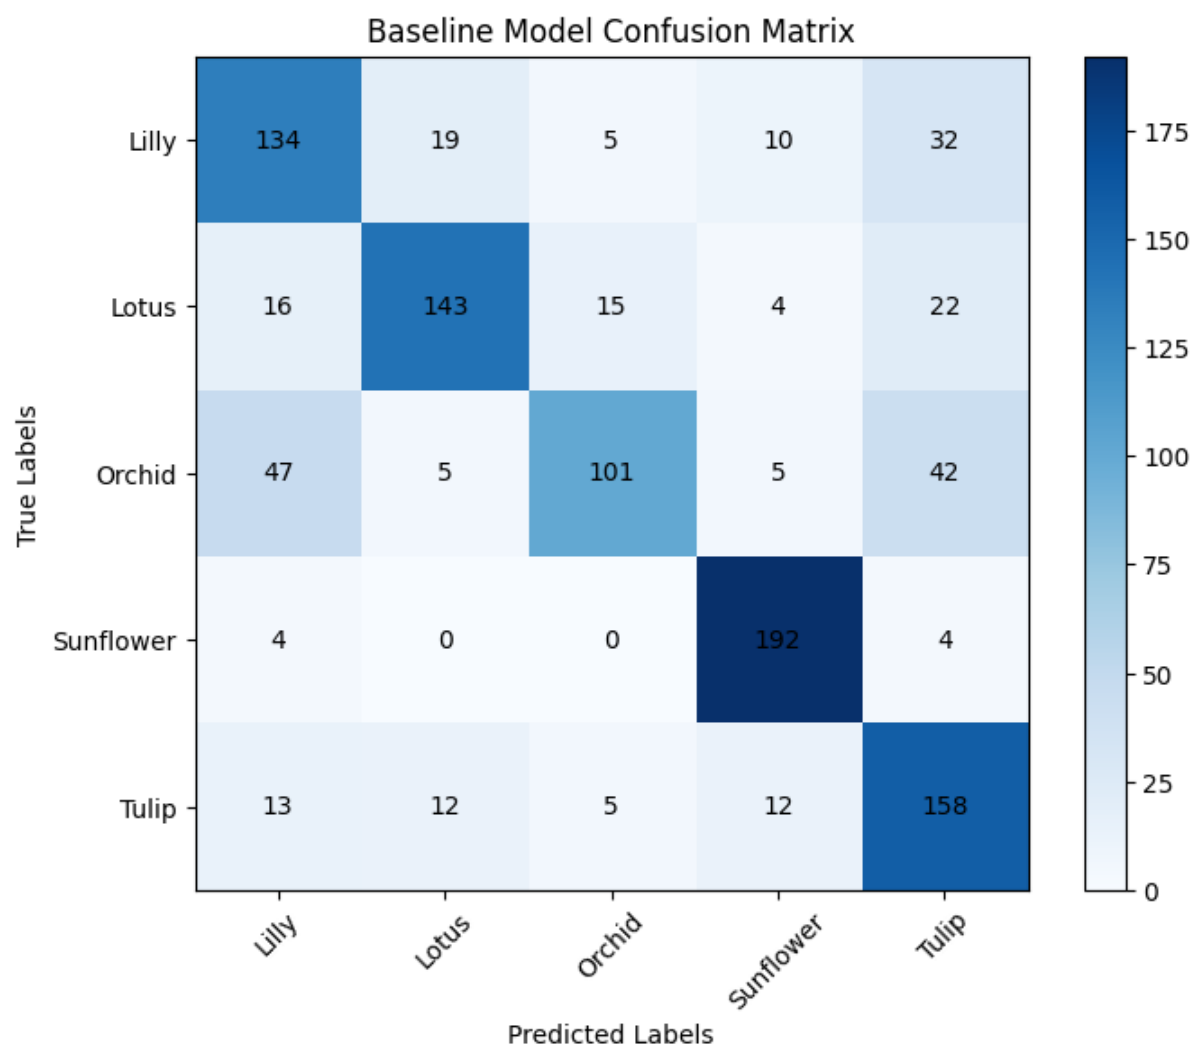
```
32/32 [=====] - 3s 96ms/step
32/32 [=====] - 5s 148ms/step
32/32 [=====] - 5s 122ms/step
```

```
In [42]: 1 print(f"Baseline Accuracy: {accuracy_score(baseline_prediction, y_true)}")
2         print(f"Resnet50 V2 Accuracy: {accuracy_score(resnet_prediction, y_true)}")
3         print(f"EfficientNetB0 Accuracy: {accuracy_score(efficientnet_prediction, y_true)}")
```


```
Baseline Accuracy: 0.728
Resnet50 V2 Accuracy: 0.874
EfficientNetB0 Accuracy: 0.876
```

```
In [48]: 1 def plot_confusion_matrix(y_true, predictions, class_names, title):
2
3     cm = confusion_matrix(y_true, predictions)
4     plt.figure(figsize=(8, 6))
5     heatmap = plt.imshow(cm, cmap='Blues')
6
7     # Set axis labels and title
8     plt.xlabel('Predicted Labels')
9     plt.ylabel('True Labels')
10    plt.title(title)
11
12    # Set xticks and yticks with class names
13    tick_labels = class_names
14    plt.xticks(ticks=np.arange(len(class_names)),
15              labels=tick_labels, rotation=45)
16    plt.yticks(ticks=np.arange(len(class_names)),
17              labels=tick_labels)
18
19    # Add numbers to the heatmap cells
20    for i in range(len(class_names)):
21        for j in range(len(class_names)):
22            plt.text(j, i, str(cm[i, j]),
23                    ha='center', va='center', color='black')
24
25    plt.colorbar(heatmap)
26    plt.show()
```


```
In [49]: 1 plot_confusion_matrix(y_true, baseline_prediction,
2                               class_names,
3                               title="Baseline Model Confusion Matrix")
```



```
In [50]: 1 plot_confusion_matrix(y_true, resnet_prediction,  
2                               class_names,  
3                               title="Resnet50 V2 Confusion Matrix")
```



```
In [51]: 1 plot_confusion_matrix(y_true, efficientnet_prediction,  
2                             class_names,  
3                             title="EfficientnetB0 Confusion Matrix")
```









```

In [59]: 1 def plot_random_images(model, val_data, classes):
2
3     images = []
4     labels = []
5     for _ in range(len(val_data)):
6         batch_images, batch_labels = next(val_data)
7         images.extend(batch_images)
8         labels.extend(batch_labels)
9
10    # Shuffle the images and labels together
11    combined = list(zip(images, labels))
12    random.shuffle(combined)
13    images, labels = zip(*combined)
14    labels = np.argmax(labels, axis=1)
15    plt.figure(figsize=(10, 6))
16    for i in range(6):
17        ax = plt.subplot(2, 3, i + 1)
18        rand_index = random.choice(range(len(images)))
19        target_image = images[rand_index]
20        pred_probs = model.predict(tf.expand_dims(target_image, axis=0), verbose=0)
21        pred_label = classes[pred_probs.argmax()]
22        true_label = classes[labels[rand_index]]
23
24        plt.imshow(target_image)
25
26        if pred_label == true_label:
27            color = "green"
28        else:
29            color = "red"
30
31        plt.title("Pred: {} {:.2f}% (True: {})".format(pred_label,
32                                                       100 * tf.reduce_max(pred_probs,
33                                                       true_label),
34                                                       color=color, fontsize=10))
35
36    plt.tight_layout()
37    plt.show()

```

```
In [60]: 1 plot_random_images(efficientnet_model, test_dataset, class_names)
```



Thank You 🙌

```
In [ ]:
```

```
1
```